

# Forecasting in Energy Markets

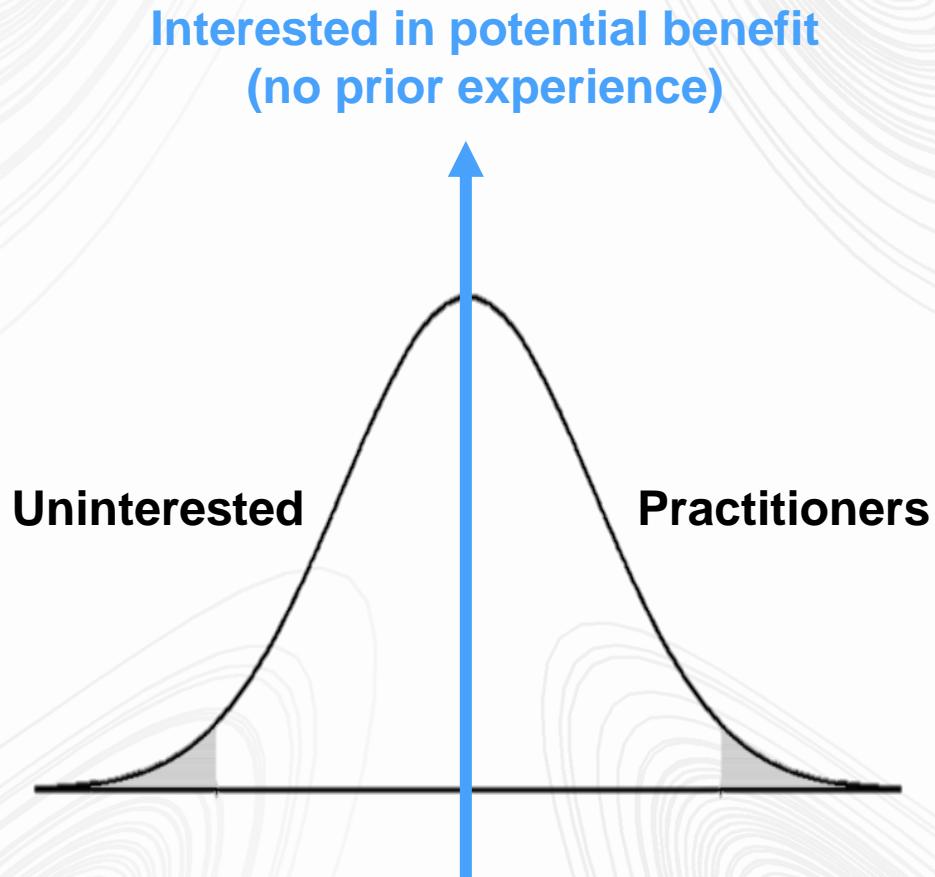
*Approaches, tools and resources for aspiring analysts*



Josh Danial  
February 2022  
NYU SPS Center for Global Affairs

# Context

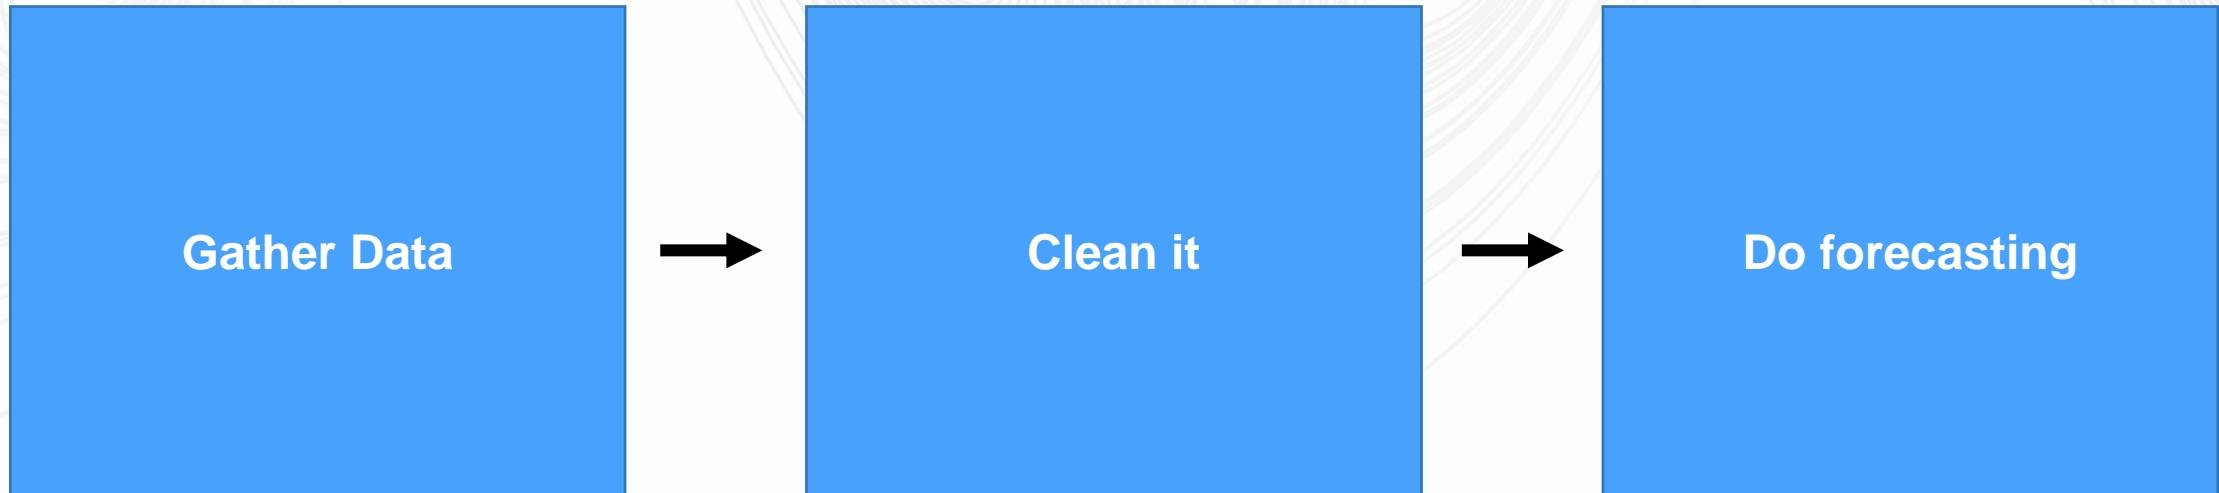
The average attendee



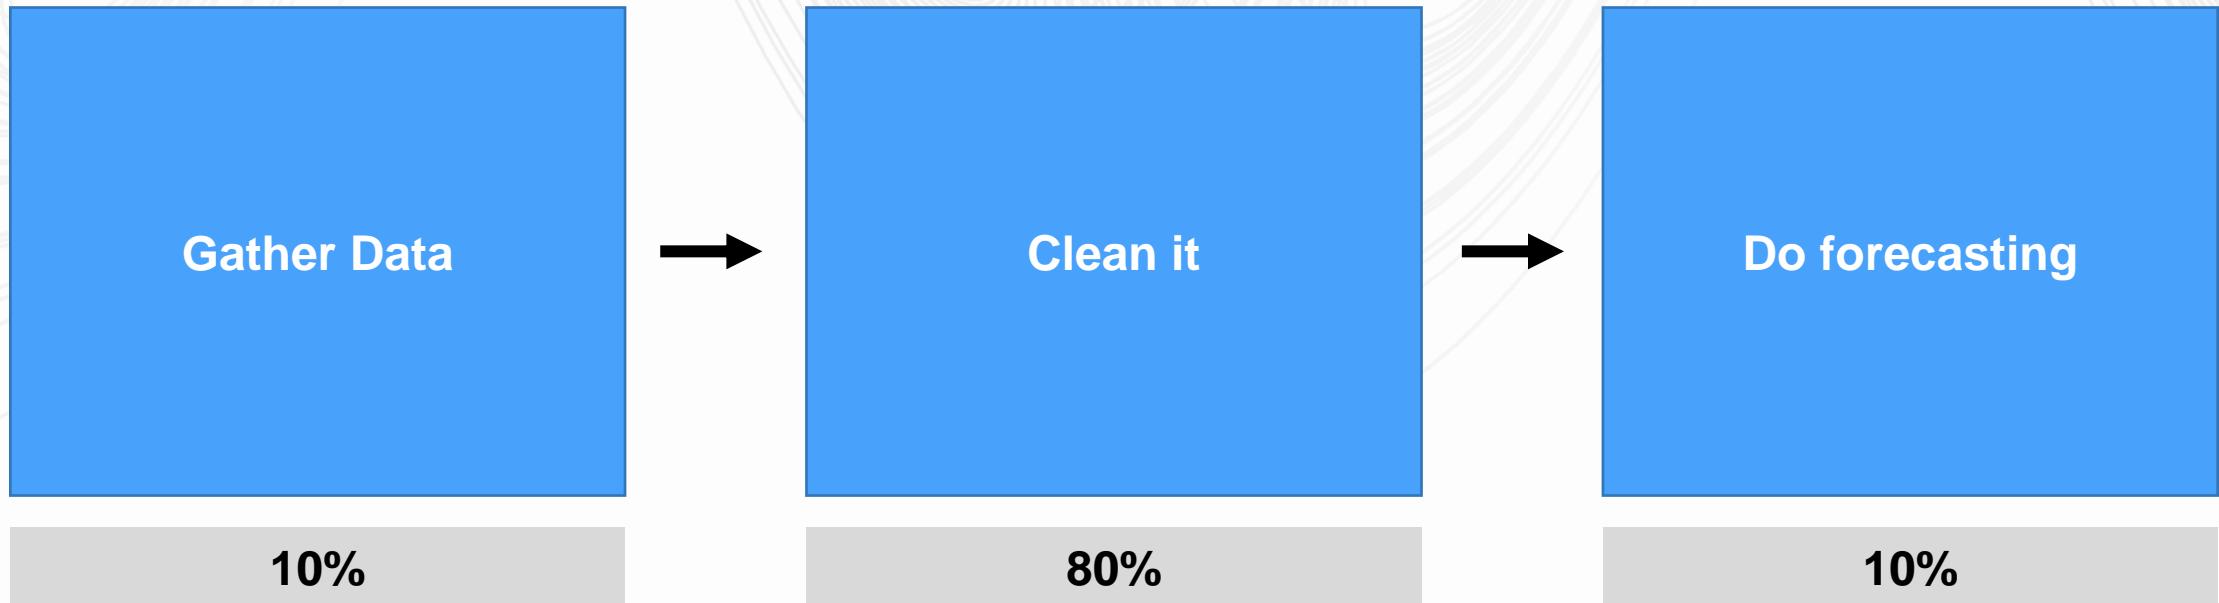
The goal: Get our feet wet



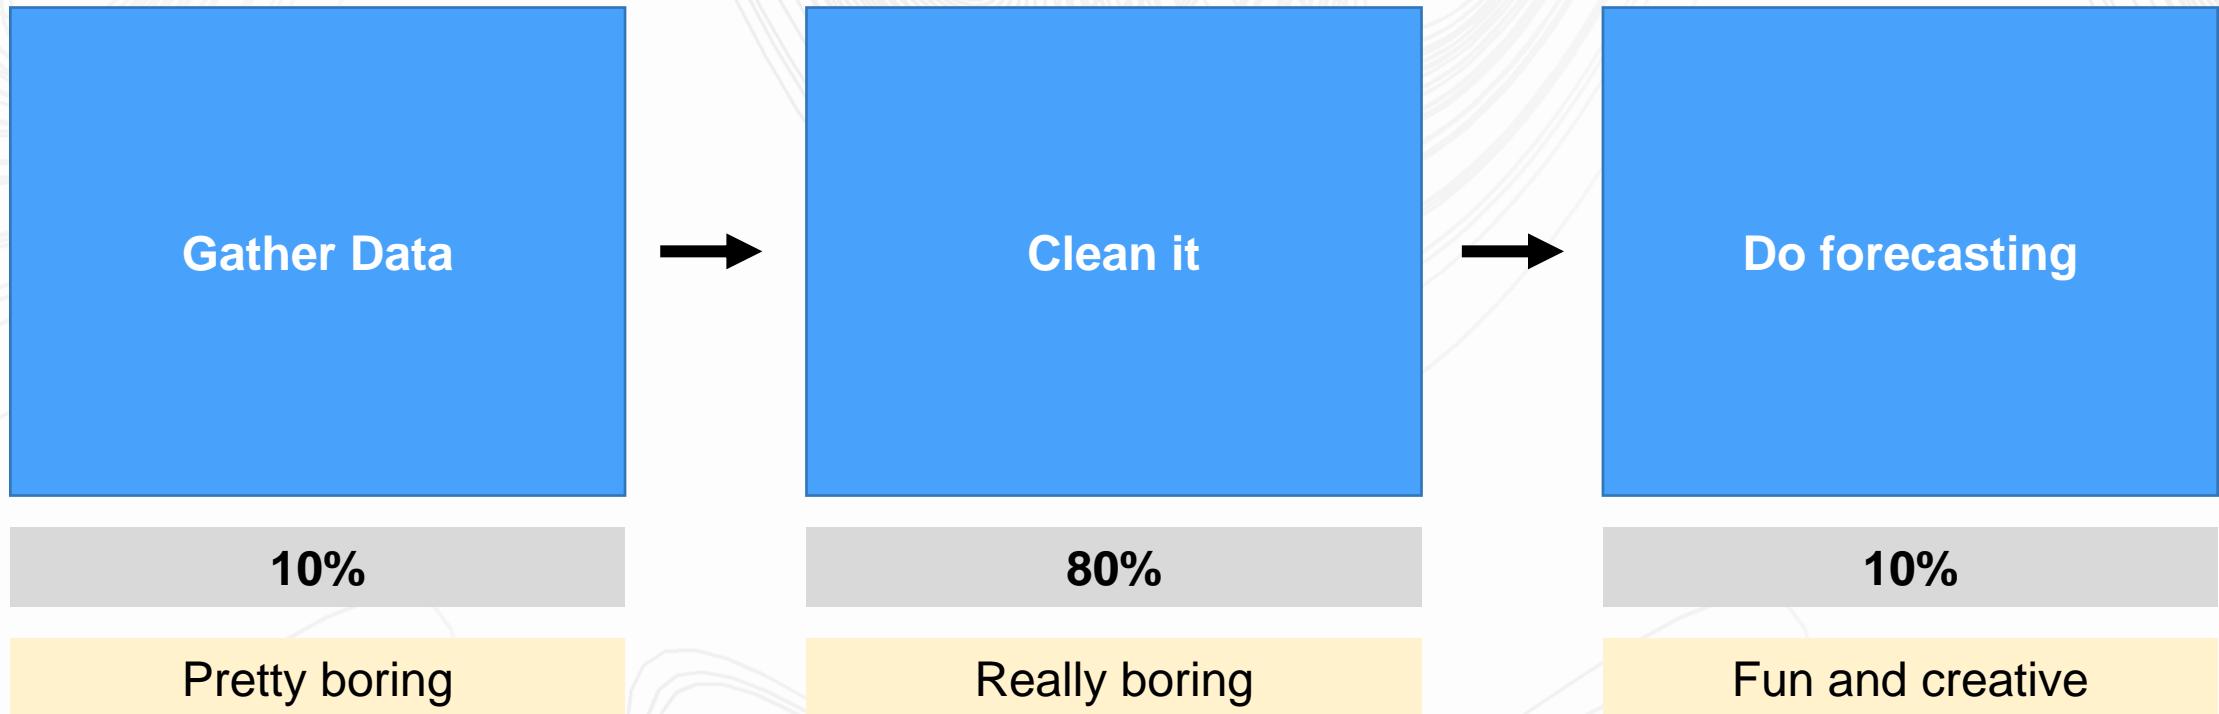
# Typical analyst workflow



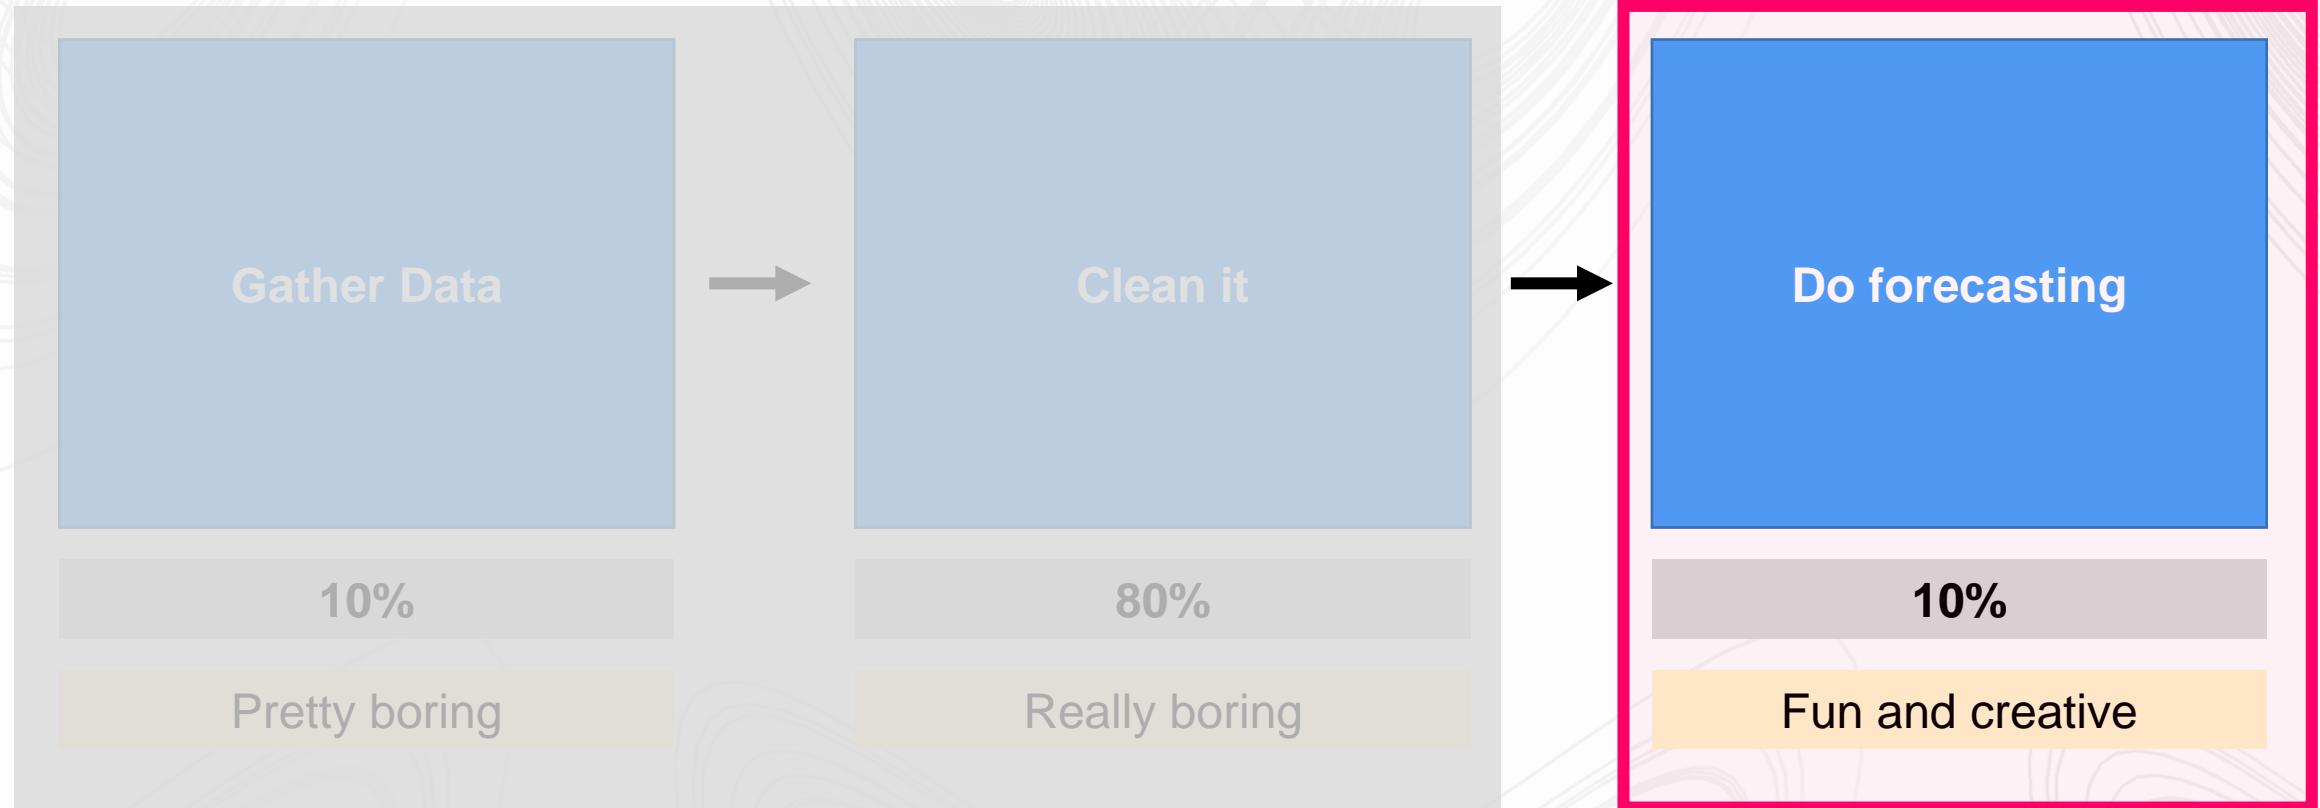
# Time allocation



# Excitement



# This talk is about the fun part



# How this relates to machine learning

For those interested

## Machine Learning

### Unsupervised

*Trained on “unlabeled” data*

### Supervised

*Trained on “labeled” data*

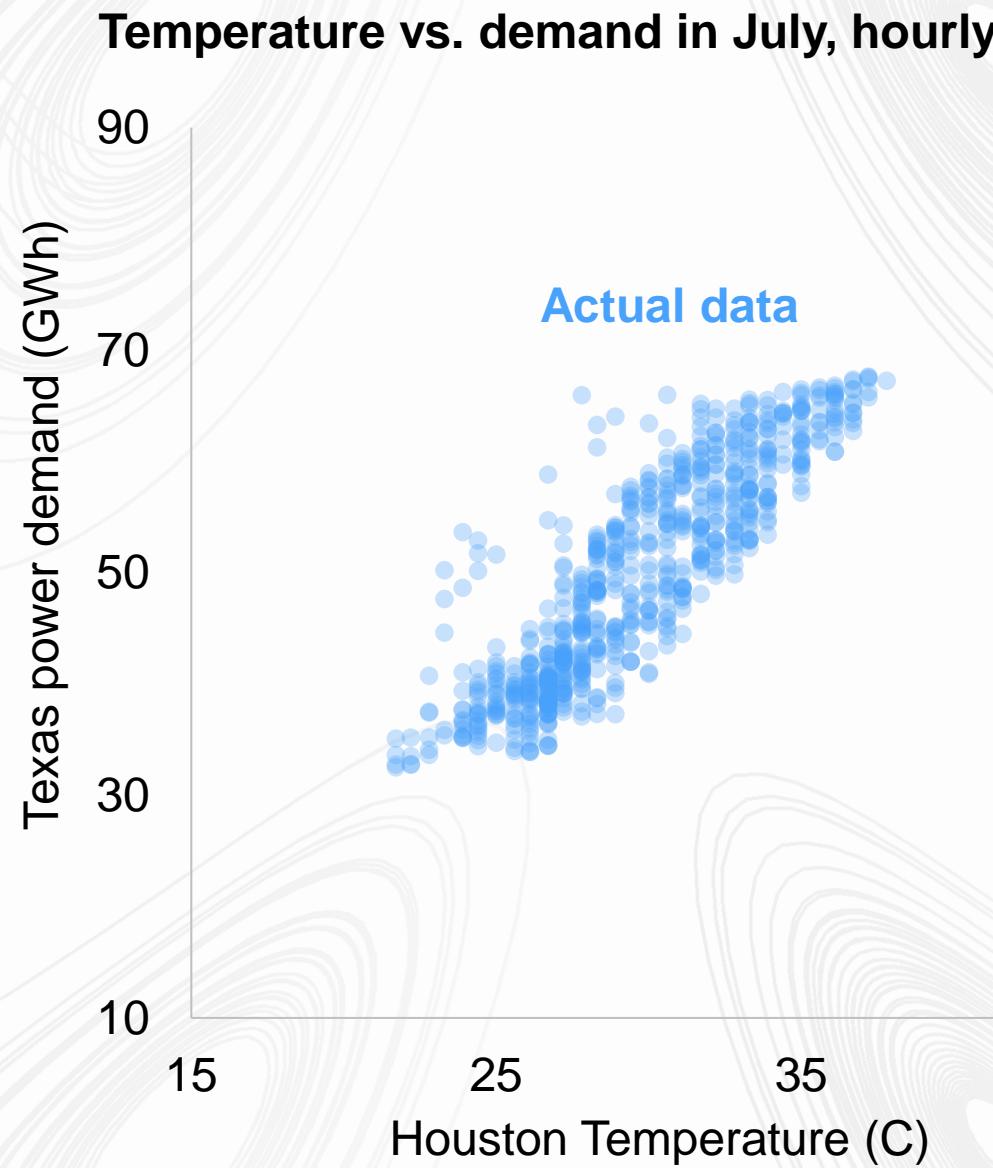
#### Classification

*Categorical prediction*

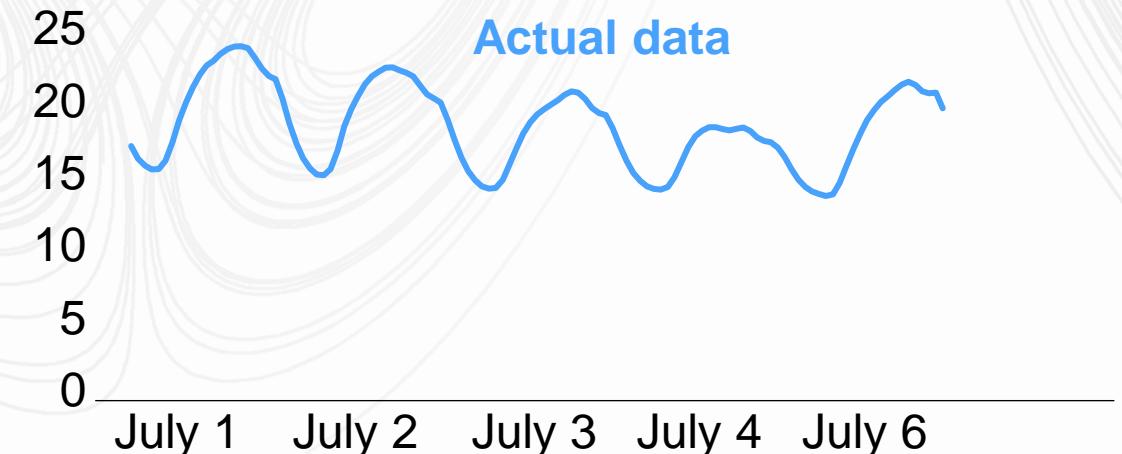
#### Regression

*Numerical prediction*

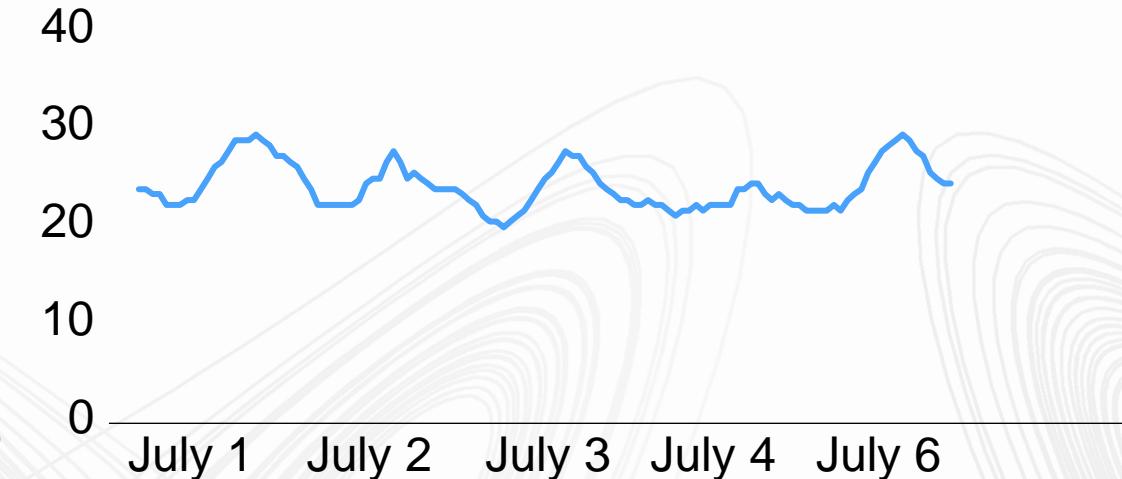
# Regressions



Power demand (GWh), hourly

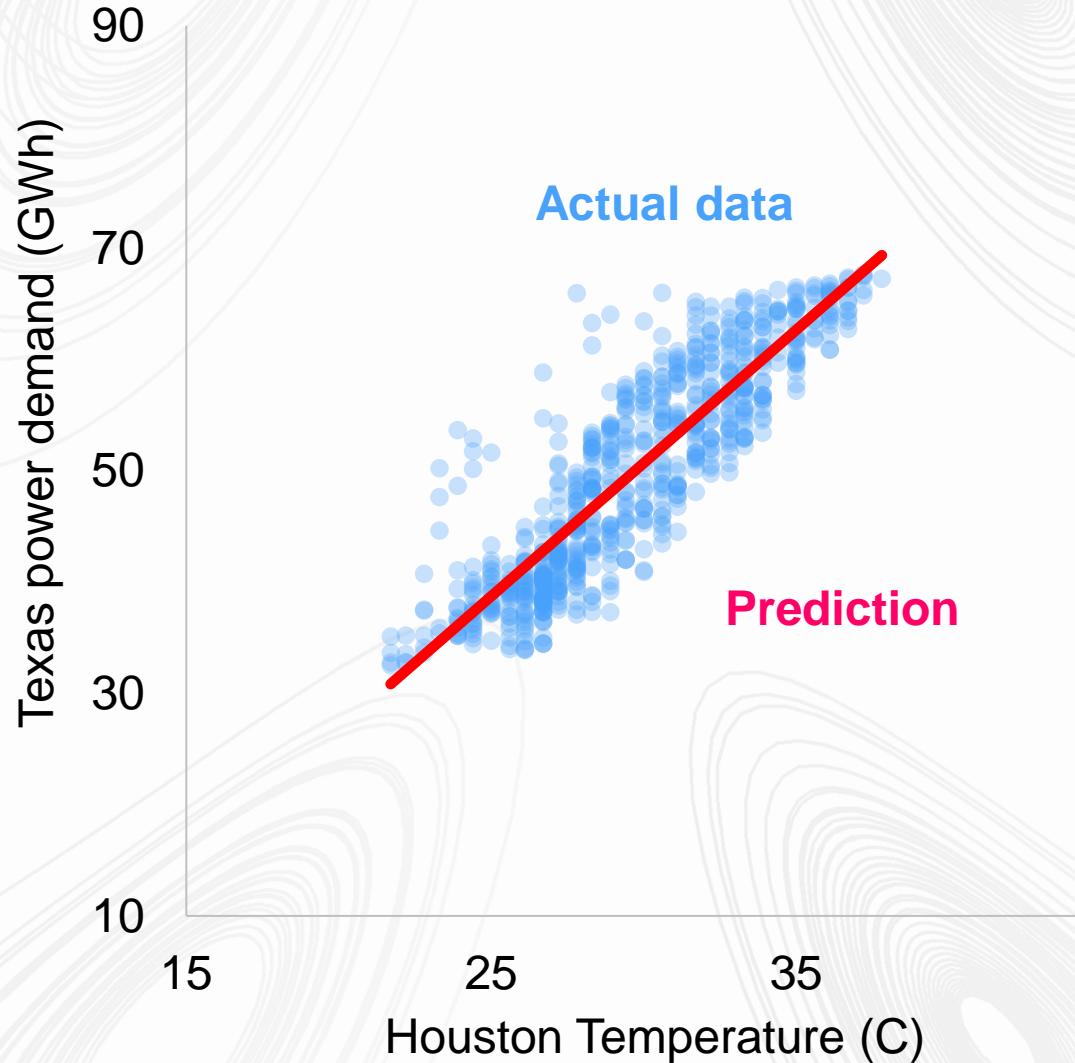


Temperature (C), hourly

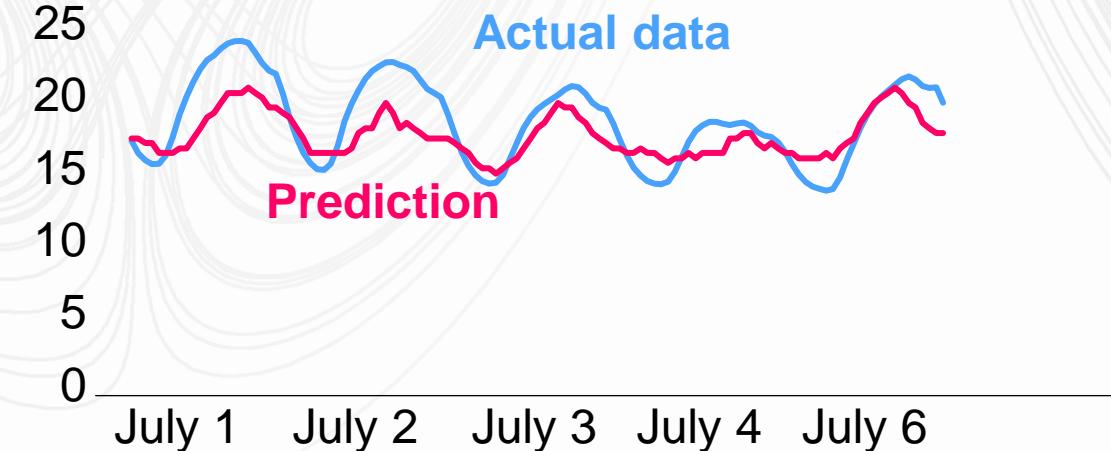


# Hunting for correlation

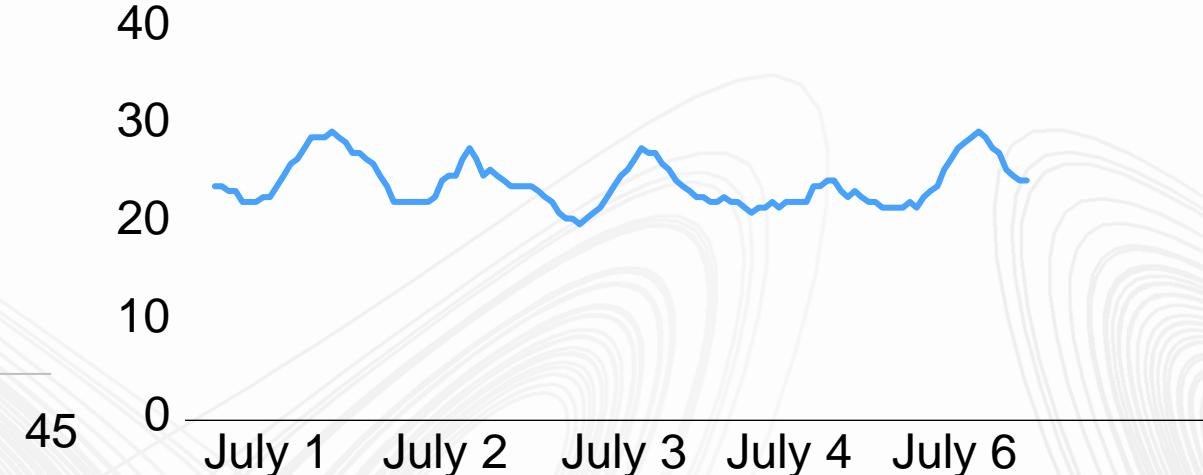
Temperature vs. demand in July, hourly



Power demand (GWh), hourly



Temperature (C), hourly



# Three misconceptions:

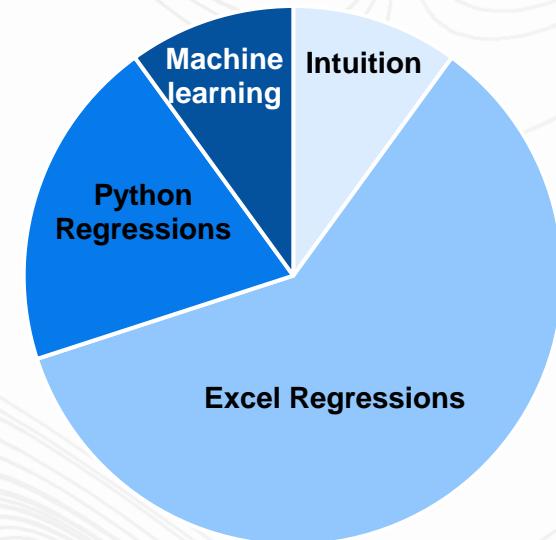
I can't...

## 1. Get the data I need

- All the data you need is publicly available
- Even at energy trading shops, most of the data feeding analysis is public

## 2. Access the best tools

Forecasting tools used in power and gas industry (my take)



## 3. Apply those techniques

- Most practitioners don't know the math that goes into a linear regression, even after using it every day
- All you need is a basic understanding of "black box testing"
  - [Wikipedia](#)
  - [Great lecture](#)

# 1. Can't get the data I need

## Energy

- [EIA open data](#)
  - Consolidated EIA data inventory
- [EIA 860 \(monthly\)](#)
  - The U.S. plant stack. Information about every power plant >1 MW.
- [EIA 923](#)
  - "Stack series." Monthly generation and fuel burn. Maps to EIA 860.
- [EIA 930](#)
  - Balancing authority level generation by tech + power demand, hourly.
- [EIA 861](#)
  - Monthly utility load and revenues. Retail power prices.
- EIA [weekly/monthly](#) petroleum supply estimates
  - Production, storage, imports, exports, refinery outputs (demand).
- [EIA drilling productivity report](#)
  - U.S. basin-level well activity by completion status, monthly.
- [JODI](#)
  - Global, monthly oil production and demand

## Socio-economic

- [FRED](#)
  - One of the easiest and most comprehensive databases of US data.
- [Bureau of Labor Statistics](#)
  - Wages, consumer price index (inflation, and much more.
- [Census](#)
  - Population/demographics— critical for a lot of hypotheses.
- [World Bank](#)
  - I don't really use this, but a lot of people do. I think most of the data is fake
- [Google's Community Mobility Reports](#)
  - Google provides county-level, global, mobility indices. Its an incredible dataset for tracking the severity of lockdowns and recoveries.
- [JHU CSSE COVID-19 Data](#)
  - County-level confirmed cases, deaths and population.

## Other

- [Data.gov](#)
  - A government project to consolidate data from a bunch of federal, state and local governments. Pretty sweet, and new datasets are added perpetually. I found a lot of the prior datasets on here.
- [Yahoo stock prices](#)
  - Stock prices
- [Fundamental data](#)
  - Consolidated, clean data quarterly data from public companies
- [VMT](#)
  - U.S. vehicle miles traveled by state and road-type, monthly.

## 2. Three Python packages, endless potential

### Data wrangling: [pandas](#)

- Nicknamed “Excel on steroids”
- Does Excel-like operations on data
- 90% of my coding time involves this package
- Resources
  - [Data Analysis with Python](#)
  - [Complete Python Pandas Data Science Tutorial](#)
  - [Web Scraping with Python Pandas](#)
  - [Web Scraping with Python](#)

### Regression: [statsmodels + scikit-learn](#)

- Statistical/machine learning regressions
- statsmodels: basic statistical models
  - [OLS](#): ordinary regression
- scikit-learn: machine learning models
  - [HuberRegressor](#): Linear regression model that is robust to outliers
  - [LinearSVR](#): Very fancy regression
  - [AdaBoostRegressor](#): Very fancy regression
  - [GradientBoostingRegressor](#): Very fancy regression

### Auto-regression: [prophet](#)

- Easy to use/automatic forecasting
- Looks at seasonality and trend in timeseries data to make a prediction
- What to use when you don't have explanatory data
  - [Quick start](#)
  - [Time series prediction using Prophet in Python](#)
  - [Forecasting gas and electricity utilization using Facebook prophet](#)

### **3. Study how people apply these tools**

#### **Some resources with methodologies and code**

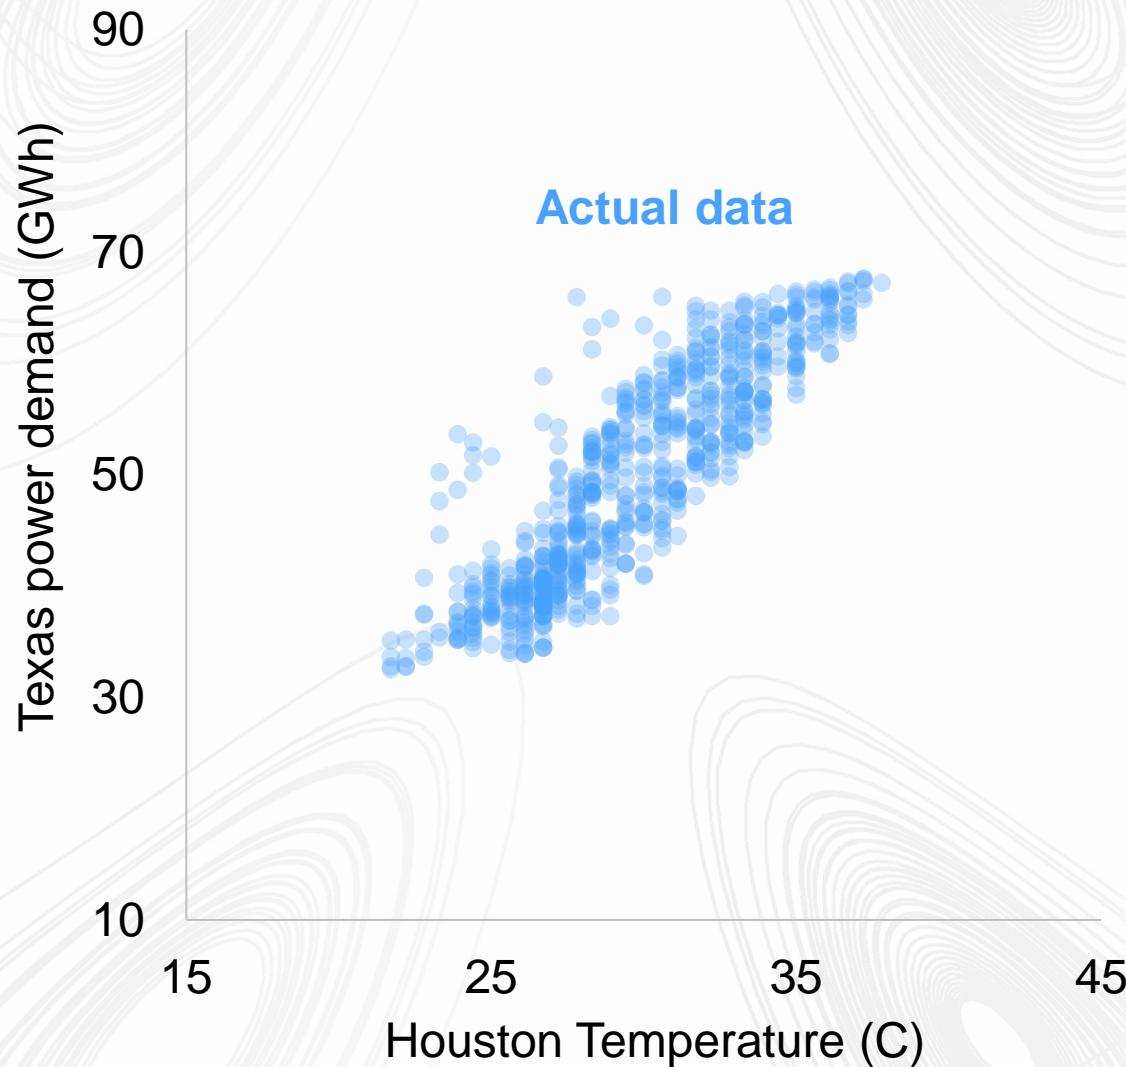
- Electricity demand
  - [Electricity demand forecasting using machine learning](#)
  - [Forecasting electricity demand with Python](#)
- Renewable generation
  - [Wind Energy Forecasting with Python](#)
  - [Forecasting Renewable Energy Generation with Streamlit and sktime](#)
- Power prices
  - [Data-driven modeling for long-term electricity price forecasting](#)
  - [Forecasting day-ahead electricity prices](#)

# **3. Code adventure**

Let's see how far we can get with our current knowledge base,  
and a couple more details

# Hunting for correlation

Temperature vs. demand in July, hourly

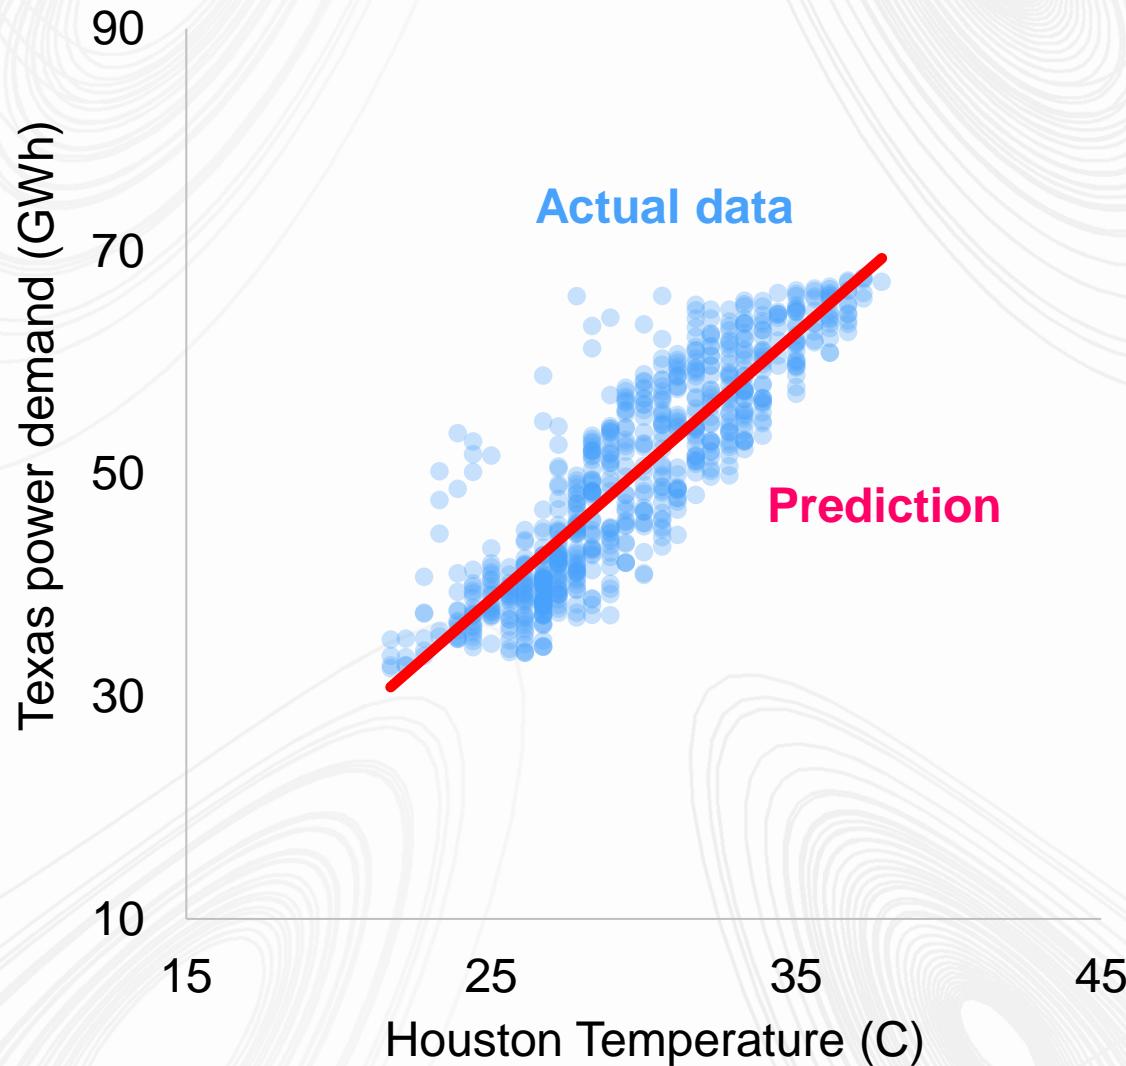


Load data/make a scatter plot

```
import pandas as pd  
  
# read  
  
data = pd.read_csv('data.csv')  
  
# chart  
  
data.plot.scatter('temp', 'load')  
  
print(data.head(5))  
  
'''Out:  
  
      datetime    load    temp  
0  2015-07-01 01:00:00  37456.0  22.8  
1  2015-07-01 02:00:00  35119.0  21.7  
2  2015-07-01 03:00:00  33638.0  21.7  
3  2015-07-01 04:00:00  32798.0  21.7  
4  2015-07-01 05:00:00  32805.0  22.2  
'''
```

# Simple, powerful linear regression

Temperature vs. demand in July, hourly



Running a regression

```
import statsmodels.api as sm

# add constant
data['constant'] = 1

# make variables
x = data[['temp', 'constant']]
y = data['load']

# statsmodels takes in Y then X
model = sm.OLS(y, x)

out = model.fit()

# apply regression function
data['pred'] = out.predict(x)

# chart
ax = data.plot.scatter('temp', 'load')

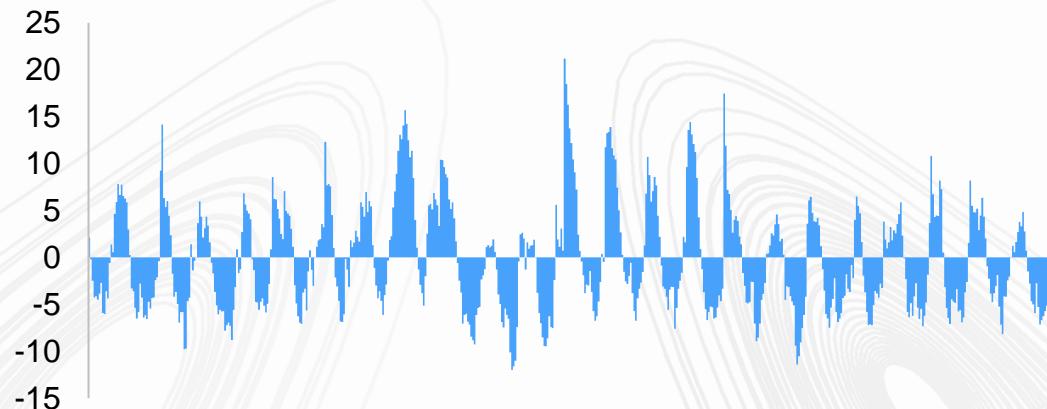
data.plot(ax=ax, x='temp', y='pred', color='r')
```

# What to look for

Actual power demand vs. expected, GWh



Errors (actual - predicted), GWh



Getting regression statistics

```
print(result.summary())
```

...Out:

OLS Regression Results

```
=====
```

Dep. Variable:	load	R-squared:	0.297
Model:	OLS	Adj. R-squared:	0.297
Method:	Least Squares	F-statistic:	2.430e+04
Date:	Thu, 03 Feb 2022	Prob (F-statistic):	0.00
Time:	05:10:55	Log-Likelihood:	-5.9964e+05
No. Observations:	57565	AIC:	1.199e+06
Df Residuals:	57563	BIC:	1.199e+06
Df Model:	1		

```
=====
```

Covariance Type: nonrobust

```
=====
```

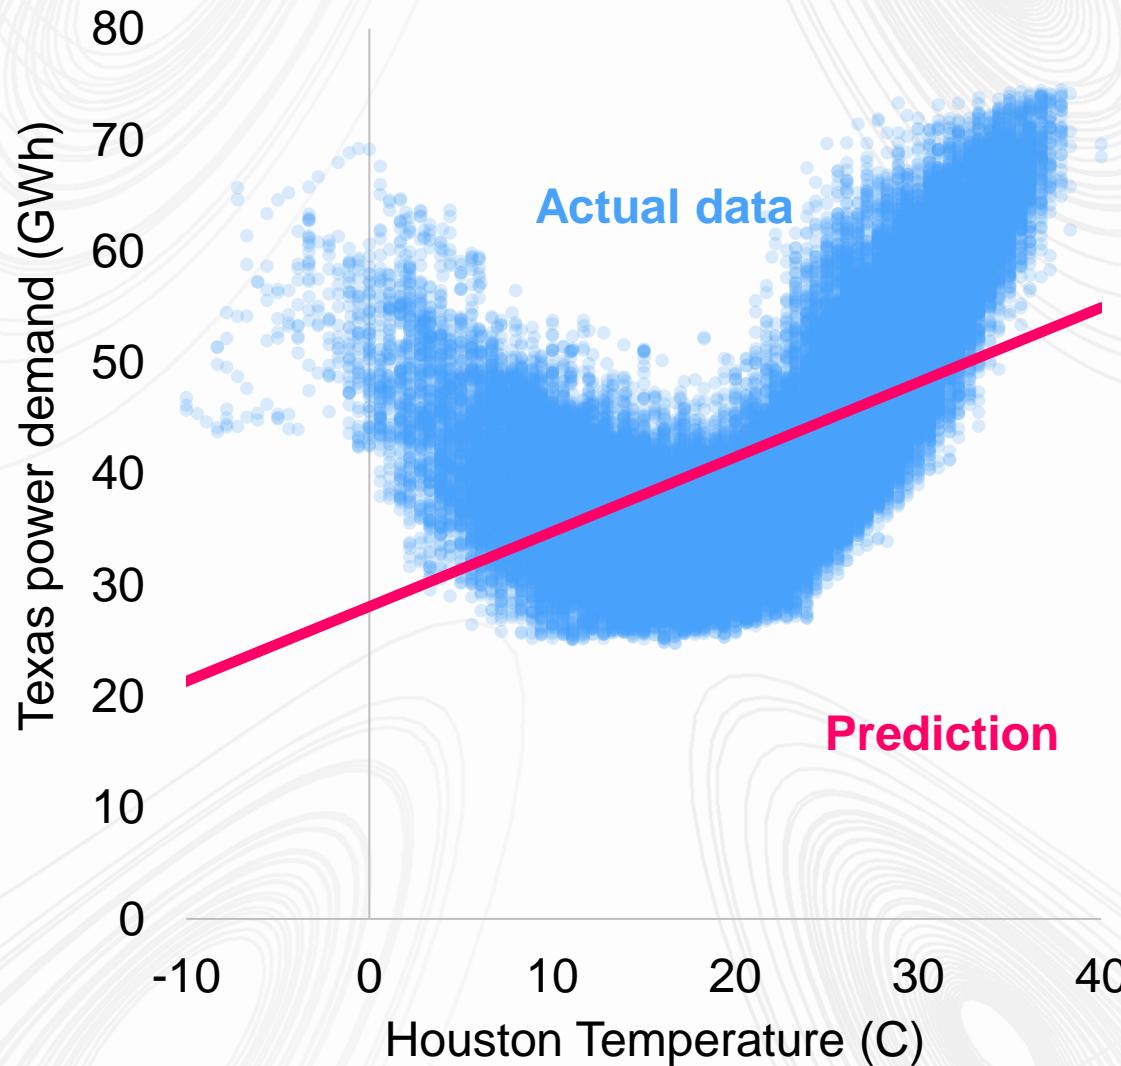
	coef	std err	t	p> t	[0.025	0.975]
temp	671.7196	4.309	155.871	0.000	663.273	680.166
constant	2.809e+04	98.890	284.066	0.000	2.79e+04	2.83e+04

```
=====
```

...

# When it breaks down

Temperature vs. demand, hourly (2015-2021)

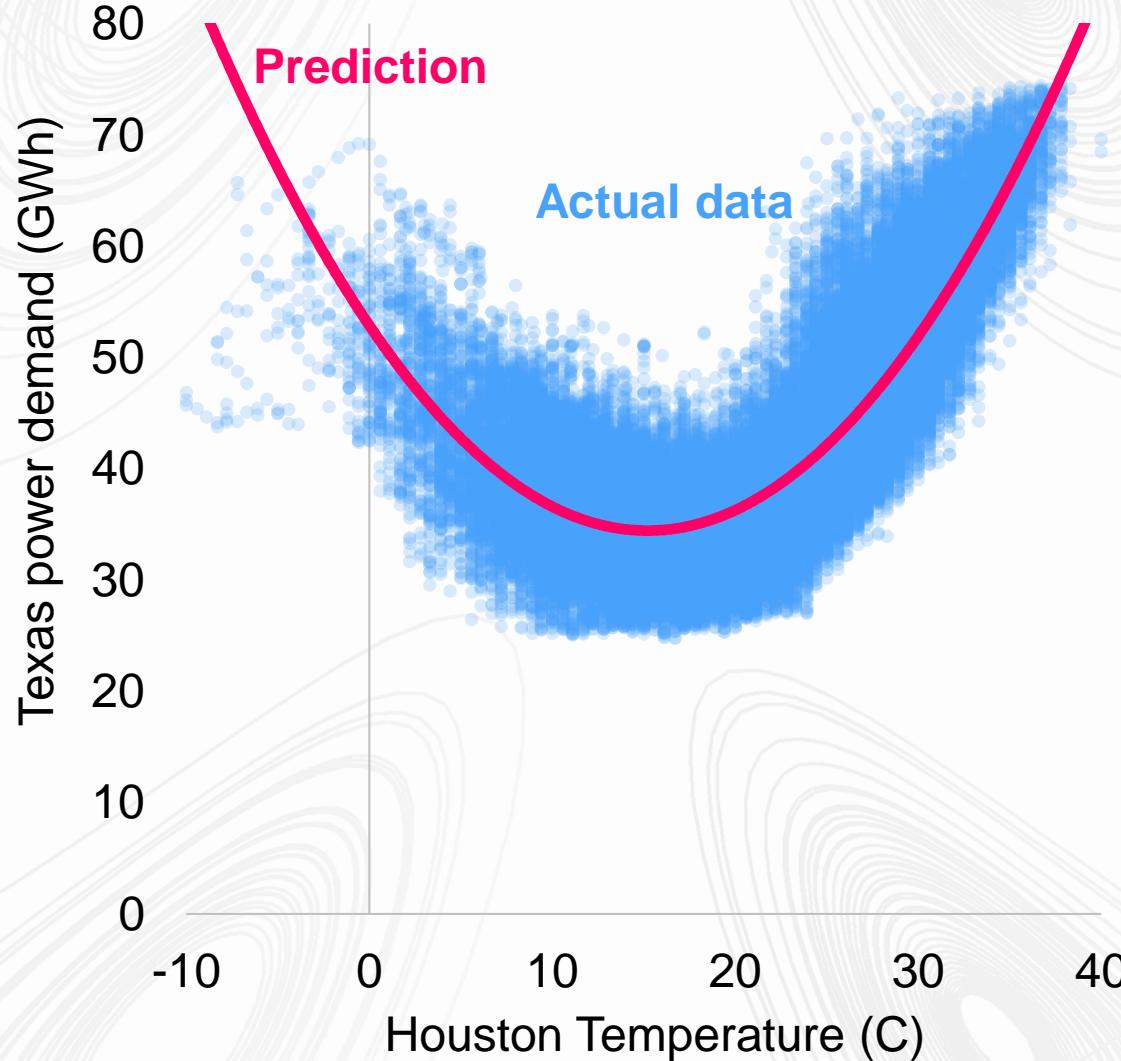


Same code from earlier

```
x = data[['temp', 'constant']]  
y = data['load']  
# same code from before  
model = sm.OLS(y, x)  
out = model.fit()  
data['pred'] = out.predict(x)  
# plot.. same code  
ax = data.plot.scatter('temp', 'load')  
data.plot(ax=ax, x='temp', y='pred', color='r')
```

# Classical statistics still has an answer

Temperature vs. demand, hourly (2015-2021)

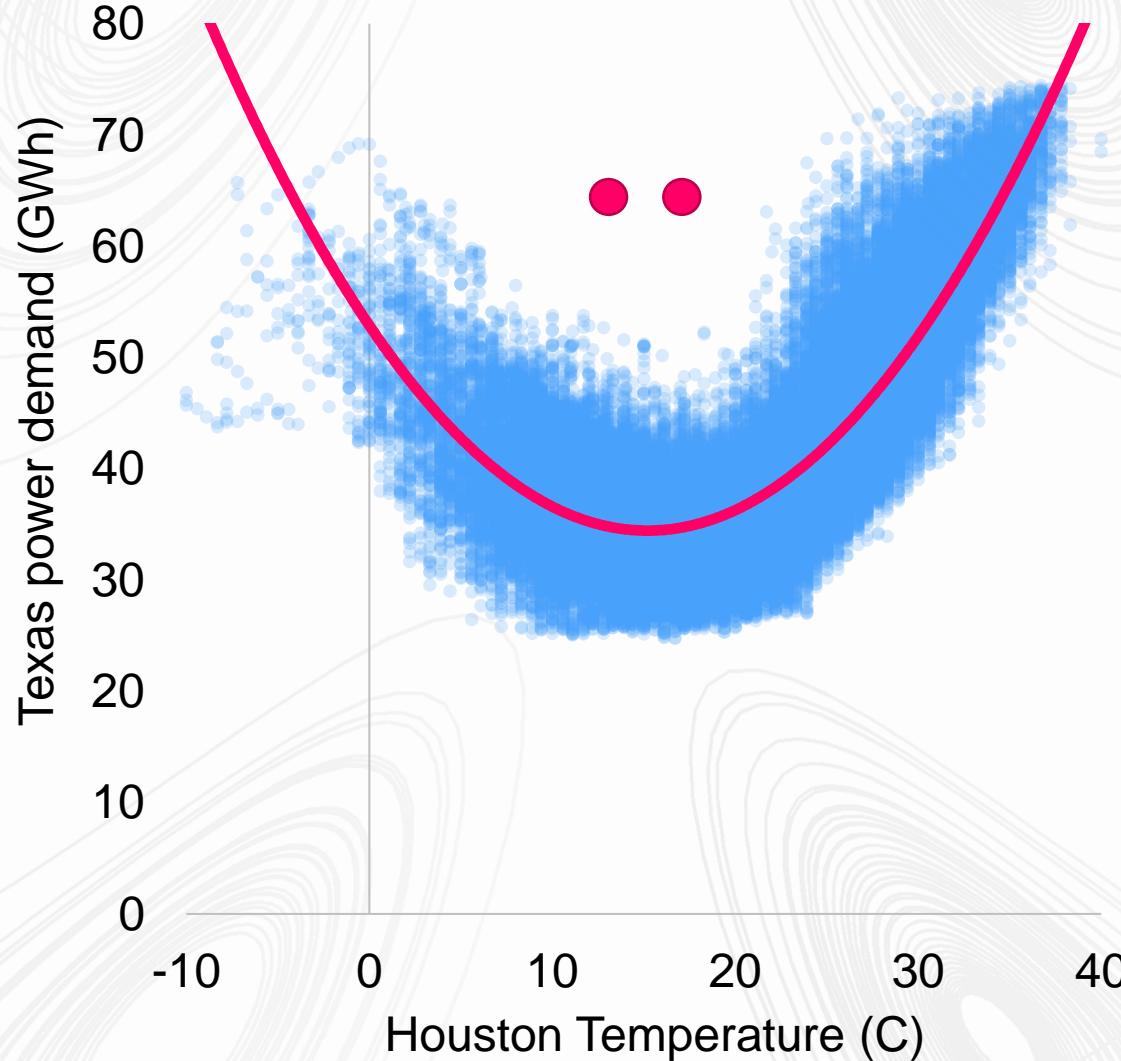


Polynomial regression in Python

```
# add polynomial  
data['temp^2'] = data['temp'] ** 2  
  
# new temp^2 variable  
x = data[['temp^2', 'temp', 'constant']]  
y = data['load']  
  
# same code from before  
model = sm.OLS(y, x)  
out = model.fit()  
data['pred'] = out.predict(x)  
  
# plot.. same code  
ax = data.plot.scatter('temp', 'load')  
data.plot(ax=ax, x='temp', y='pred', color='r')
```

# Classical statistics still has an answer

Temperature vs. demand, hourly (2015-2021)



Polynomial regression in Python

```
# add polynomial
data['temp^2'] = data['temp'] ** 2

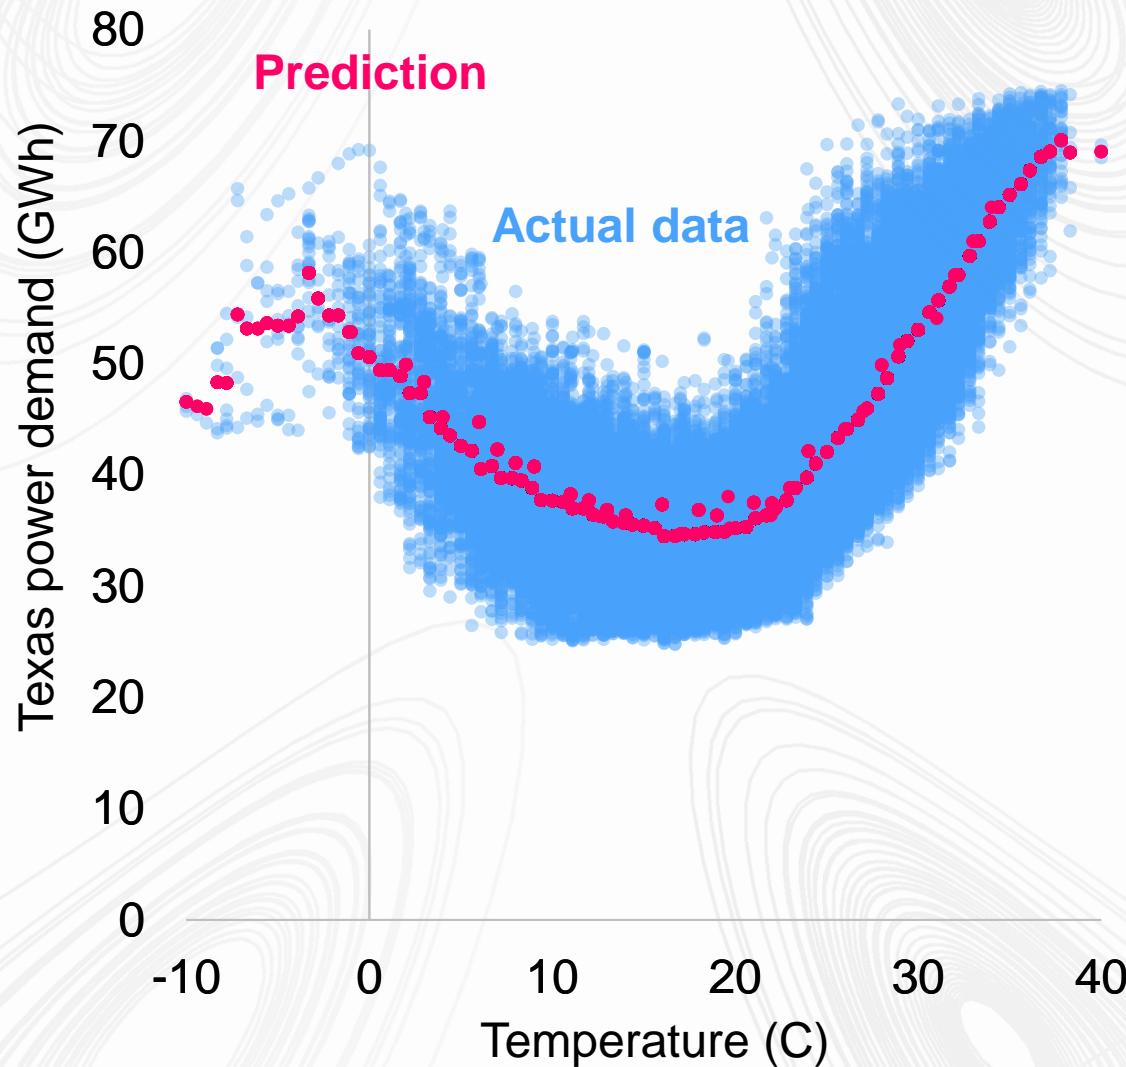
# new temp^2 variable
x = data[['temp^2', 'temp', 'constant']]
y = data['load']

# same code from before
model = sm.OLS(y, x)
out = model.fit()
data['pred'] = out.predict(x)

# plot.. same code
ax = data.plot.scatter('temp', 'load')
data.plot(ax=ax, x='temp', y='pred', color='r')
```

# Let's turn it up anyways

Temperature vs. demand, hourly (2015-2021)

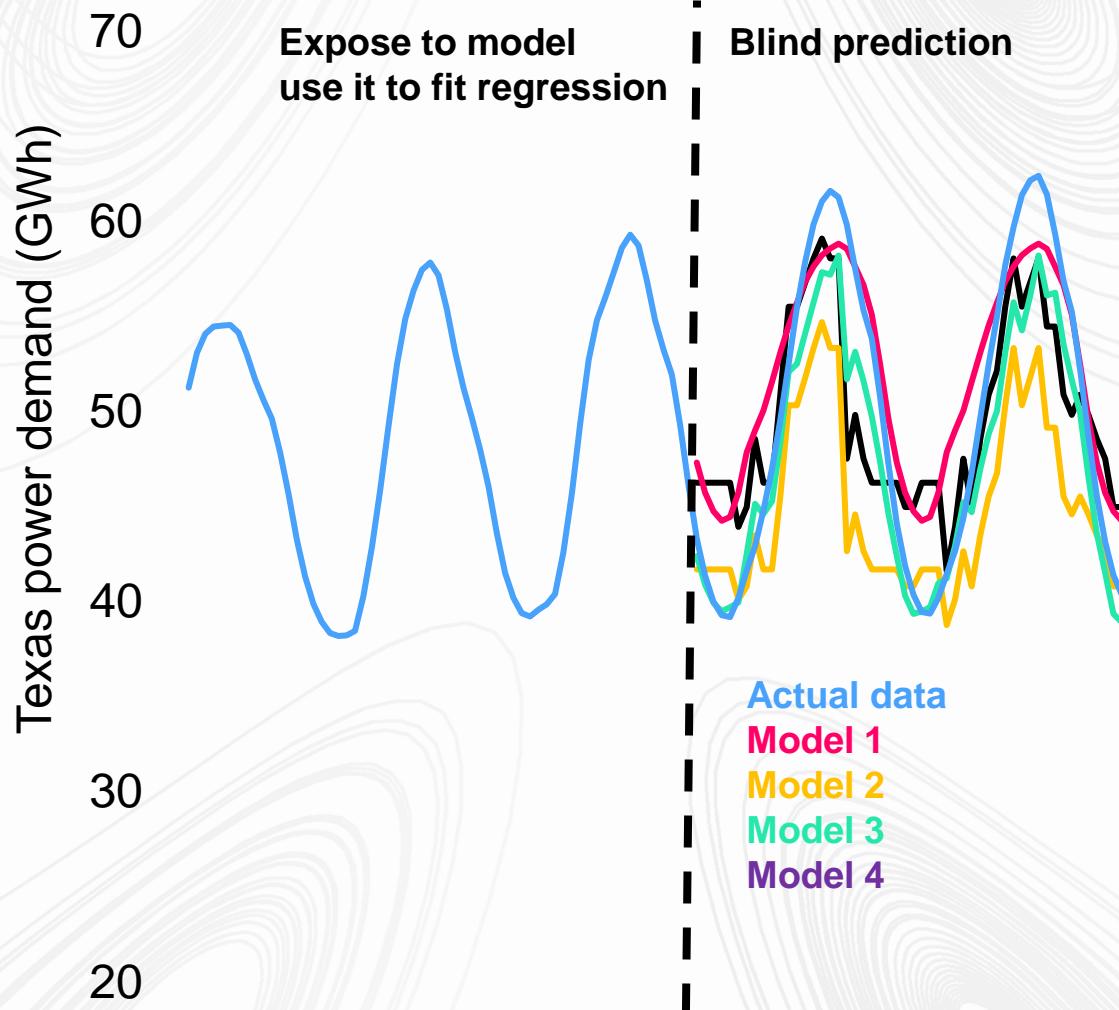


Using a machine learning regressor

```
from sklearn.ensemble import GradientBoostingRegressor  
  
# fancy machine learning  
  
model = GradientBoostingRegressor()  
  
# no constant, no polynomial  
  
x = data[['temp']]  
y = data['load']  
  
model = model.fit(x, y)  
  
data['pred'] = model.predict(x)  
  
# chart... same code  
  
ax = data.plot.scatter('temp', 'load')  
data.plot(ax=ax, x='temp', y='pred', color='r')
```

# Backtest

## Illustrative backtest



## Train/test split

```
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.ensemble import AdaBoostRegressor

xgb = GradientBoostingRegressor()
ada = AdaBoostRegressor()

data['datetime'] = pd.to_datetime(data['datetime'])

# split data
train = data[data['datetime'] < '1/1/2019']
test = data[data['datetime'] > '1/1/2019']

x_train, y_train = train[['temp']], train[['load']]
x_test, y_test = test[['temp']], test[['load']]

# fit/predict
xgb = xgb.fit(x_train, y_train)
ada = ada.fit(x_train, y_train)

y_test['xgb'] = xgb.predict(x_test)
y_test['ada'] = ada.predict(x_test)

print(y_test.head(3))

'''Out:
      load          xgb          ada
0  36953.0  39659.922366  41642.391821
1  37114.0  39252.865195  41642.391821
2  37154.0  39659.922366  41642.391821
```

# **Thank you!**